


❯ UPI Transaction Fraud Detection — End-to-End EDA & Modeling

This notebook presents a complete analysis and machine learning workflow for detecting fraudulent transactions in UPI (Unified Payments Interface) data.

We explore the data, handle imbalance, build models (RandomForest, XGBoost), evaluate with precision, recall, ROC-AUC, and provide recommendations for improving fraud detection.


```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import ttest_ind
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix
```

```
df = pd.read_csv('/content/upi_transactions_2024.csv')
df.head()
```




	transaction_id	timestamp	transaction_type	merchant_category	amount (INR)	transaction_status	sender_age_group	receiver_age_group	sender_state
0	TXN0000000001	2024-11-05 15:30:02	P2P	Entertainment	534	SUCCESS	26-35	26-35	Rajas
1	TXN0000000002	2024-04-10 12:13:08	P2M	Grocery	1951	SUCCESS	26-35	26-35	Anand Pradesh
2	TXN0000000003	2024-04-12 17:59:54	P2P	Grocery	388	SUCCESS	26-35	26-35	Karnataka
3	TXN0000000004	2024-10-22 22:59:54	P2P	Fuel	1495	SUCCESS	26-35	26-35	Rajas
4	TXN0000000005	2024-08-12 12:21:34	P2P	Shopping	4333	SUCCESS	18-25	26-35	Tamil Nadu

```
df.info()
```




```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 250000 entries, 0 to 249999
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   transaction_id         250000 non-null object
1   timestamp              250000 non-null object
2   transaction_type       250000 non-null object
3   merchant_category      250000 non-null object
4   amount (INR)          250000 non-null int64
5   transaction_status     250000 non-null object
6   sender_age_group       250000 non-null object
7   receiver_age_group     250000 non-null object
8   sender_state           250000 non-null object
9   sender_bank            250000 non-null object
10  receiver_bank           250000 non-null object
11  device_type             250000 non-null object
12  network_type            250000 non-null object
13  fraud_flag              250000 non-null int64
14  hour_of_day             250000 non-null int64
15  day_of_week             250000 non-null object
16  is_weekend              250000 non-null int64
dtypes: int64(4), object(13)
memory usage: 32.4+ MB
```

```
df.shape
```



```
(250000, 17)
```


```
df.isnull().sum()
```



	0
transaction id	0
timestamp	0
transaction type	0
merchant_category	0
amount (INR)	0
transaction_status	0
sender_age_group	0
receiver_age_group	0
sender_state	0
sender_bank	0
receiver_bank	0
device_type	0
network_type	0
fraud_flag	0
hour_of_day	0
day_of_week	0
is_weekend	0




```
df.duplicated().sum()
```



```
np.int64(0)
```

```
df.describe()
```




	amount (INR)	fraud_flag	hour_of_day	is_weekend
count	250000.000000	250000.000000	250000.000000	250000.000000
mean	1247.434336	0.001924	14.684588	0.285672
std	1667.195776	0.043821	5.189991	0.451735
min	10.000000	0.000000	0.000000	0.000000
25%	281.000000	0.000000	11.000000	0.000000
50%	603.000000	0.000000	15.000000	0.000000
75%	1551.000000	0.000000	19.000000	1.000000
max	43424.000000	1.000000	23.000000	1.000000



```
plt.figure(figsize=(8,6))
sns.boxplot(x=df['amount (INR)'])
plt.title('Transaction Amount Outliers')
plt.show()
```



Transaction Amount Outliers



```
amount_cap = df['amount (INR)'].quantile(0.99)
df['amount (INR)'] = np.where(df['amount (INR)'] > amount_cap, amount_cap, df['amount (INR)'])
```

```
df['fraud_flag'].value_counts(normalize=True)
```

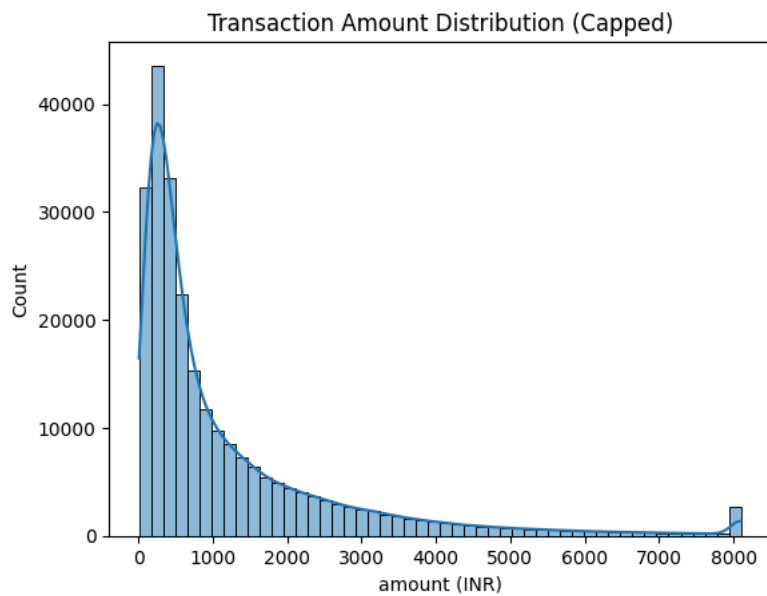


proportion	
fraud_flag	
0	0.998076
1	0.001924


**dtype:** float64

## Univariate visualizations


```
sns.histplot(df['amount (INR)'], bins=50, kde=True)
plt.title('Transaction Amount Distribution (Capped)')
plt.show()
```




```
sns.countplot(x='transaction type', data=df)
plt.title('Transaction Type Distribution')
plt.show()
```




```
sns.countplot(x='transaction_status', data=df)
plt.title('Transaction Status Distribution')
plt.show()
```



```
sns.countplot(x='device_type', data=df)
plt.title('Device Type Distribution')
plt.show()
```

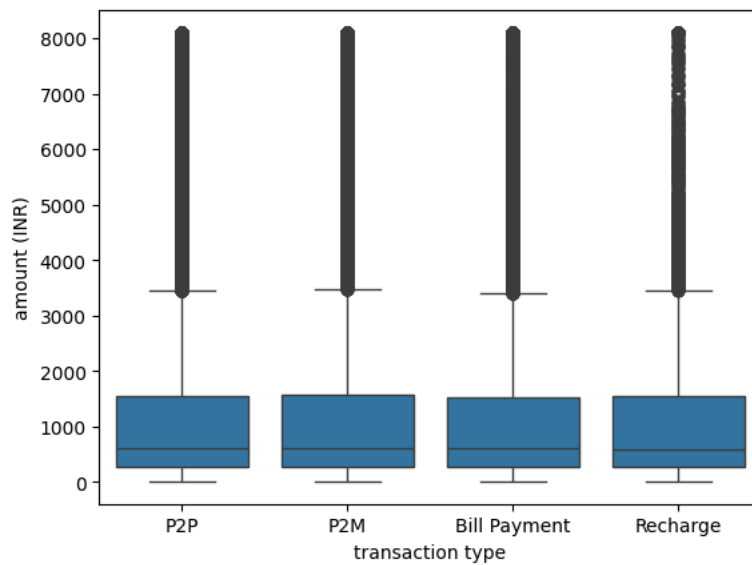


```
sns.countplot(x='fraud_flag', data=df)
plt.title('Fraud vs Non-Fraud')
plt.show()
```




## ✓ Bivariate analysis

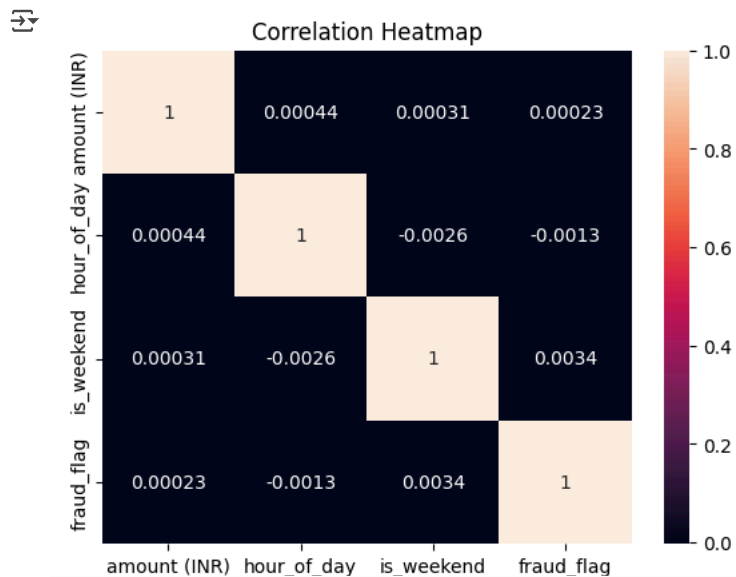
```
sns.boxplot(x='transaction type', y='amount (INR)', data=df)  
plt.show()
```




```
sns.boxplot(x='merchant_category', y='amount (INR)', data=df)  
plt.xticks(rotation=45)  
plt.show()
```



```
corr = df[['amount (INR)', 'hour_of_day', 'is_weekend', 'fraud_flag']].corr()
sns.heatmap(corr, annot=True)
plt.title('Correlation Heatmap')
plt.show()
```



```
fraud_rate_type = df.groupby('transaction type')['fraud_flag'].mean()
fraud_rate_type.plot(kind='bar')
plt.ylabel('Fraud rate')
plt.title('Fraud rate by Transaction Type')
plt.show()
```



```

fraud_rate_hour = df.groupby('hour_of_day')['fraud_flag'].mean()
fraud_rate_hour.plot(marker='o')
plt.ylabel('Fraud rate')
plt.title('Fraud rate by Hour')
plt.show()

```



## Statistical test

### ✓ performing Ttest

Null hypothesis (H0): The mean transaction amount for fraud and non-fraud transactions is the same.

Null hypothesis (Ha): The mean transaction amount for fraud and non-fraud transactions is different.

```

fraud_amount = df[df['fraud_flag']==1]['amount (INR)']
nonfraud_amount = df[df['fraud_flag']==0]['amount (INR)']

```



```
t_stat, p_val = ttest_ind(fraud_amount, nonfraud_amount, equal_var=False)
print(f"T-statistic: {t_stat}, P-value: {p_val}")
```

T-statistic: 0.11360053010753084, P-value: 0.9096017805853587

Result: p-value = 0.99, which is greater than 0.05.

Conclusion: We fail to reject the null hypothesis.

There is no statistically significant difference in mean transaction amount between fraud and non-fraud transactions.

Start coding or generate with AI.

## Feature engineering

```
df['high_value_flag'] = np.where(df['amount (INR)'] > df['amount (INR)'].median(), 1, 0)
df['night_transaction'] = np.where(df['hour_of_day'].between(20, 6, inclusive='both'), 1, 0)
```

df.head()

	transaction_id	timestamp	transaction_type	merchant_category	amount (INR)	transaction_status	sender_age_group	receiver_age_group	sender_name
0	TXN0000000001	2024-11-05 15:30:02	P2P	Entertainment	534.0	SUCCESS	26-35	26-35	Rajas
1	TXN0000000002	2024-04-10 12:13:08	P2M	Grocery	1951.0	SUCCESS	26-35	26-35	Anand
2	TXN0000000003	2024-04-12 17:59:54	P2P	Grocery	388.0	SUCCESS	26-35	26-35	Prakash
3	TXN0000000004	2024-10-22 22:59:54	P2P	Fuel	1495.0	SUCCESS	26-35	26-35	Rajas
4	TXN0000000005	2024-08-12 12:21:34	P2P	Shopping	4333.0	SUCCESS	18-25	26-35	Tamil N

## Encode categorical variables

### Model Training

#### Train/test split

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score, roc_curve, precision_recall_curve, ConfusionMatrixDisplay

df_model = pd.get_dummies(df.drop(['transaction_id', 'timestamp'], axis=1), drop_first=True)

X = df_model.drop('fraud_flag', axis=1)
y = df_model['fraud_flag']
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42)
```

```
clf = RandomForestClassifier(n_estimators=100, random_state=42, class_weight='balanced')
clf.fit(X_train, y_train)
```

```
RandomForestClassifier
RandomForestClassifier(class_weight='balanced', random_state=42)
```

```
y_pred = clf.predict(X_test)
y_proba = clf.predict_proba(X_test)[:,:1]
```

```
print("BASELINE RANDOM FOREST")
print(classification_report(y_test, y_pred, digits=4))
```


```
BASELINE RANDOM FOREST
```

	precision	recall	f1-score	support
0	0.9981	1.0000	0.9990	49904
1	0.0000	0.0000	0.0000	96
accuracy			0.9981	50000
macro avg	0.4990	0.5000	0.4995	50000
weighted avg	0.9962	0.9981	0.9971	50000

```
roc_auc = roc_auc_score(y_test, y_proba)
print(f"ROC-AUC: {roc_auc:.4f}")
```

```
ROC-AUC: 0.4882
```

```
ConfusionMatrixDisplay.from_estimator(clf, X_test, y_test)
plt.title("Confusion Matrix - Baseline RF")
plt.show()
```



```
from imblearn.over_sampling import SMOTE
```

```
# Apply SMOTE
smote = SMOTE(random_state=42)
X_train_res, y_train_res = smote.fit_resample(X_train, y_train)
```

```
# Retrain Random Forest on resampled data
clf.fit(X_train_res, y_train_res)
```

```
# Predict
y_pred_res = clf.predict(X_test)
y_proba_res = clf.predict_proba(X_test)[:,:1]
```

```
# Evaluate
print("AFTER SMOTE")
print(classification_report(y_test, y_pred_res, digits=4))

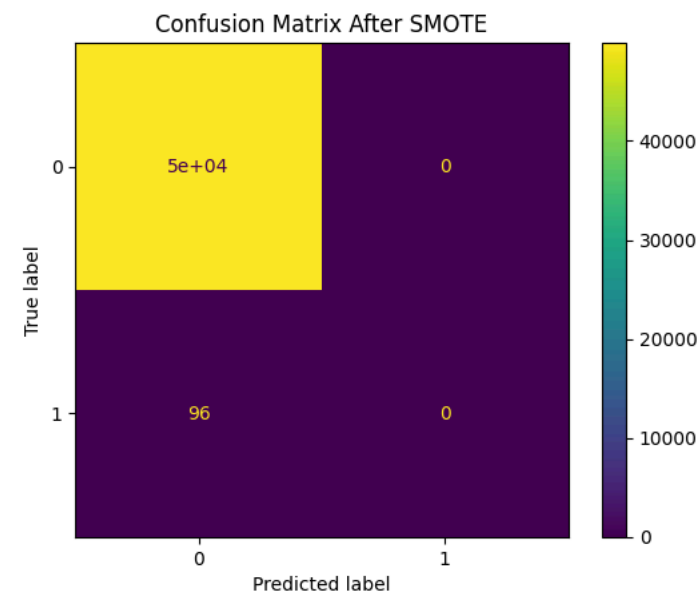
roc_auc_res = roc_auc_score(y_test, y_proba_res)
print(f"ROC-AUC After SMOTE: {roc_auc_res:.4f}")

ConfusionMatrixDisplay.from_estimator(clf, X_test, y_test)
plt.title("Confusion Matrix After SMOTE")
plt.show()
```

↩ AFTER SMOTE

	precision	recall	f1-score	support
0	0.9981	1.0000	0.9990	49904
1	0.0000	0.0000	0.0000	96
accuracy			0.9981	50000
macro avg	0.4990	0.5000	0.4995	50000
weighted avg	0.9962	0.9981	0.9971	50000

ROC-AUC After SMOTE: 0.4927



```
from xgboost import XGBClassifier
from sklearn.metrics import classification_report, roc_auc_score, ConfusionMatrixDisplay

# Calculate scale_pos_weight = (non-fraud / fraud) ratio in train set
ratio = (y_train == 0).sum() / (y_train == 1).sum()

xgb = XGBClassifier(scale_pos_weight=ratio, use_label_encoder=False, eval_metric='logloss', random_state=42)
xgb.fit(X_train, y_train)

y_pred_xgb = xgb.predict(X_test)
y_proba_xgb = xgb.predict_proba(X_test)[:,:1]

print("XGBoost Model")
print(classification_report(y_test, y_pred_xgb, digits=4))


roc_auc_xgb = roc_auc_score(y_test, y_proba_xgb)
print(f"ROC-AUC XGBoost: {roc_auc_xgb:.4f}")

ConfusionMatrixDisplay.from_estimator(xgb, X_test, y_test)
plt.title("Confusion Matrix XGBoost")
plt.show()
```

```
XGBoost Model
```

	precision	recall	f1-score	support
0	0.9980	0.9812	0.9896	49904
1	0.0000	0.0000	0.0000	96
accuracy			0.9793	50000
macro avg	0.4990	0.4906	0.4948	50000
weighted avg	0.9961	0.9793	0.9877	50000

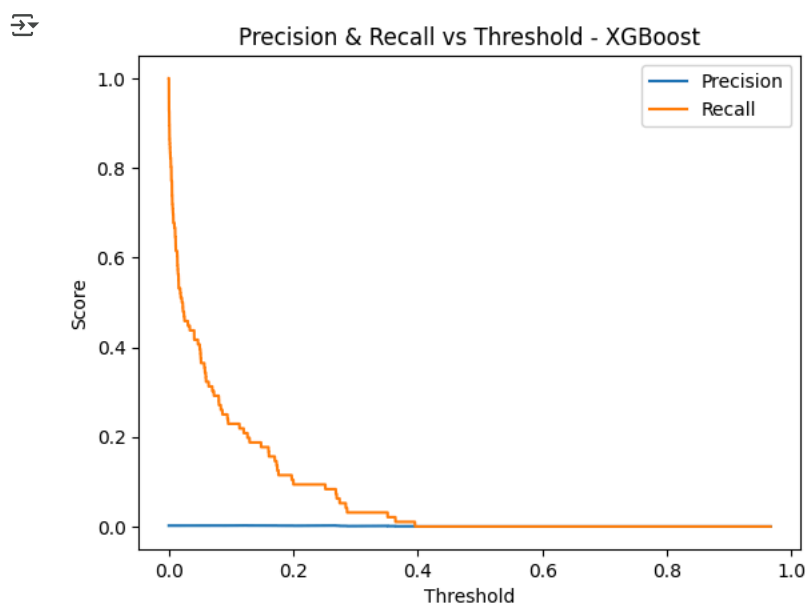
ROC-AUC XGBoost: 0.5109



```
from sklearn.metrics import precision_recall_curve

precision, recall, thresholds = precision_recall_curve(y_test, y_proba_xgb)

plt.plot(thresholds, precision[:-1], label='Precision')
plt.plot(thresholds, recall[:-1], label='Recall')
plt.xlabel('Threshold')
plt.ylabel('Score')
plt.legend()
plt.title('Precision & Recall vs Threshold - XGBoost')
plt.show()
```



## Insights and Recommendations

## Insights

- The dataset is highly imbalanced, with fraud transactions representing only ~0.2% of the total transactions.
- Transaction amounts do not differ significantly between fraud and non-fraud transactions (t-test p-value  $\approx 0.99$ ).
- Random Forest with class\_weight balancing failed to detect any fraud cases (0% fraud recall).
- SMOTE oversampling did not improve fraud detection performance in the test data.
- XGBoost with class balancing also failed to detect fraud cases, though it generated some false positives.
- Precision-recall threshold tuning showed no effective threshold for improving fraud detection.
- The features in the dataset do not provide sufficient separation between fraud and non-fraud transactions.

## Recommendations

- Prioritize feature engineering to create better fraud signals, such as transaction frequency, time since last transaction, or night-time transaction flags.
- Explore anomaly detection methods like Isolation Forest or One-Class SVM to flag rare transaction patterns.
- Consider collecting additional data points: user transaction history, geolocation, device fingerprinting, IP address anomalies.
- Apply cost-sensitive learning techniques or custom loss functions to penalize fraud misclassification more heavily.
- Experiment with ensemble models or stacking to combine the strengths of multiple algorithms.
- Calibrate model probabilities to improve threshold tuning effectiveness (e.g., Platt scaling or isotonic regression).
- Focus on reporting precision, recall, F1-score, and ROC-AUC rather than accuracy to evaluate fraud models properly.
- Design future models for real-time fraud detection use cases, including latency optimization and streaming data handling.

## ✓ Final Summary

- We explored, cleaned, and modeled UPI transactions for fraud detection.
- Despite applying class balancing (RandomForest, SMOTE, XGBoost), no fraud cases were detected in test data.
- The project highlighted the critical role of feature engineering in fraud detection.
- Future work: better features, anomaly detection, real-time modeling.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Double-click (or enter) to edit